CSEN 296B-2 | Week 2.1 Lab

Working Agreement for AstraNotes

David Nguyen | April 6, 2026

1. Working Agreement

This agreement defines how I will manage my AstraNotes project as a solo developer using an AI-native workflow for the remainder of the quarter.

Planning and Tracking Work

I will maintain a lightweight backlog in a GitHub Project board with columns for Backlog, In Progress, Review, and Done. Each task will be scoped small enough to finish in one working session. At the start of each week I will pick two to three items from the backlog that align with the current course milestone and move them into In Progress. I will not pull new work until the current items clear my Definition of Done.

How I Will Use AI Inside My Workflow

AI, primarily Claude, will serve as a drafting and reasoning partner. I will use it for generating boilerplate code, exploring design alternatives, writing first-draft documentation, and rubber-ducking architectural decisions. AI will not be treated as a source of truth. Every output that touches AstraNotes source code, requirements, or documentation must pass through my own review before it is committed or submitted.

Documenting Prompts, Decisions, and Revisions

For any significant design or implementation decision, I will keep a short entry in an Architecture Decision Log stored in the repo. Each entry will note the context, what I asked AI, what it suggested, and what I actually chose along with why. Minor code-generation prompts do not need full logs, but any prompt that shapes a feature boundary, data model change, or security choice will be recorded.

Deciding Whether AI Output Is Acceptable

Before accepting AI-generated output I will ask three questions. First, can I explain what this code or text does without re-reading the prompt? Second, does it fit the existing AstraNotes architecture, specifically the Next.js App Router conventions, Supabase Row Level Security policies, and the component patterns already in use? Third, does it introduce any dependency, scope change, or assumption I did not ask for? If any answer is no, I will revise or reject the output.

Preventing Drift, Duplication, and Low-Quality Output

I will enforce a one-feature-at-a-time rule. Feature branches will be short-lived and merged only after passing my Definition of Done. I will run the existing lint and type-check scripts before every commit. If I notice AI repeatedly suggesting patterns that conflict with my current codebase conventions, I will stop and update my prompt context rather than patching the output after the fact.

2. Reflection

Writing this agreement forced me to acknowledge a pattern I have already fallen into on AstraNotes: accepting AI output that looks correct without verifying it against the project's actual file structure and conventions. In earlier labs I merged generated code that used a slightly different import path than the rest of the codebase, which caused a build error I had to debug later. Having an explicit rule to check architectural fit before accepting output would have caught that. Going forward, the three-question acceptance check and the decision log should keep my workflow disciplined even when deadlines are tight. The agreement also reminds me that planning work in small, trackable chunks is not overhead; it is what keeps a solo AI-native project from quietly drifting off course.